

Task 3:

```
class SGDUnivariateLinearRegression:

    def __init__(self):
        self.theta_0: float = 0.
        self.theta_1: float = 0.
        self.rng = np.random.default_rng(RANDOM_SEED)

    def predict(self, x):
        y = self.theta_0 + self.theta_1 * x
        return y

    def fit(self, x, y, n_iter: int = 100, learning_rate: float = 1.0):
        for t in range(n_iter):
            sample_ix = self.rng.integers(0, len(x))

            xt = x[sample_ix]
            yt = y[sample_ix]
            temp0 = self.theta_0 - learning_rate * (self.theta_0 + self.theta_1 * xt - yt)
            temp1 = self.theta_1 - learning_rate * (self.theta_0 + self.theta_1 * xt - yt) * xt
            self.theta_0 = temp0
            self.theta_1 = temp1

        return self
```

```
[83] SGDUnivariateLinearRegression = SGDUnivariateLinearRegression()
SGDUnivariateLinearRegression.fit(x, y1)
pred = SGDUnivariateLinearRegression.predict(x)
def mse(y_pred, y_true):
    squared = ((y_pred - y_true)**2).mean()

    return squared

print('y1: ', mse(pred, y1))
print('y2: ', mse(pred, y2))
```

```
y1: 0.23665898370500849
y2: 0.6367834505153138
```

```

n_iters = [50, 100, 200, 500, 1000, 2000]
learning_rates = [1., .1, .01]

# we plot the MSE achieved by the closed form model as a reference
closed_form = UnivariateLinearRegression()
closed_form.fit(x, y1)
mse_base = mse(y_pred=closed_form.predict(x), y_true=y1)
plt.plot(n_iters, np.ones_like(n_iters) * mse_base, label="closed form", linestyle='--', c='b')


for alpha in learning_rates:
    mses = []
    for n_iter in n_iters:
        SGDUnivariateLinearRegression = SGDUnivariateLinearRegression()
        SGDUnivariateLinearRegression.fit(x, y1, n_iter, alpha)
        pred = SGDUnivariateLinearRegression.predict(x)

        # compute its mse and append the mse value to the mses list

        mse_ = mse(pred, y1)
        mses.append(mse_)
    plt.plot(n_iters, mses, label=f"alpha = {alpha:.2f}")

plt.xlabel("n_iter")
plt.ylabel("MSE")
plt.legend()
plt.show()

```




alpha = 1 is too big. The other two are correct but alpha = 0.1 converges quicker.

Task 4:

```
[103] from sklearn.linear_model import LinearRegression
```

```
lng = LinearRegression()  
lng.fit(X, y)  
predx = lng.predict(X)  
plot_residuals(predx, y)
```



No basic assumptions are not fulfilled

Task 4b

The estimated parameters θ of the linear model can be found in the `.coef_` member variable. The feature names can be found in the `.feature_names_in_` member variable. They are the same as the names of the columns of `x` and should be in the same order.

Using these, answer the following questions:

- Which are the 3 most influential features?
- How do you interpret the sign of the coefficients?
- If you had to exclude 1 feature, which one would you select and why?

```
[106] lng.coef_
```

```
array([ -10.01219782, -239.81908937,  519.83978679,  324.39042769,  
       -792.18416163,  476.74583782,  101.04457032,  177.06417623,  
        751.27932109,   67.62538639])
```

```
[107] lng.feature_names_in_
```

```
array(['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6'],  
      dtype=object)
```

3 most influential = bmi, bp, s2 exclude = s1 seems to be the least influential